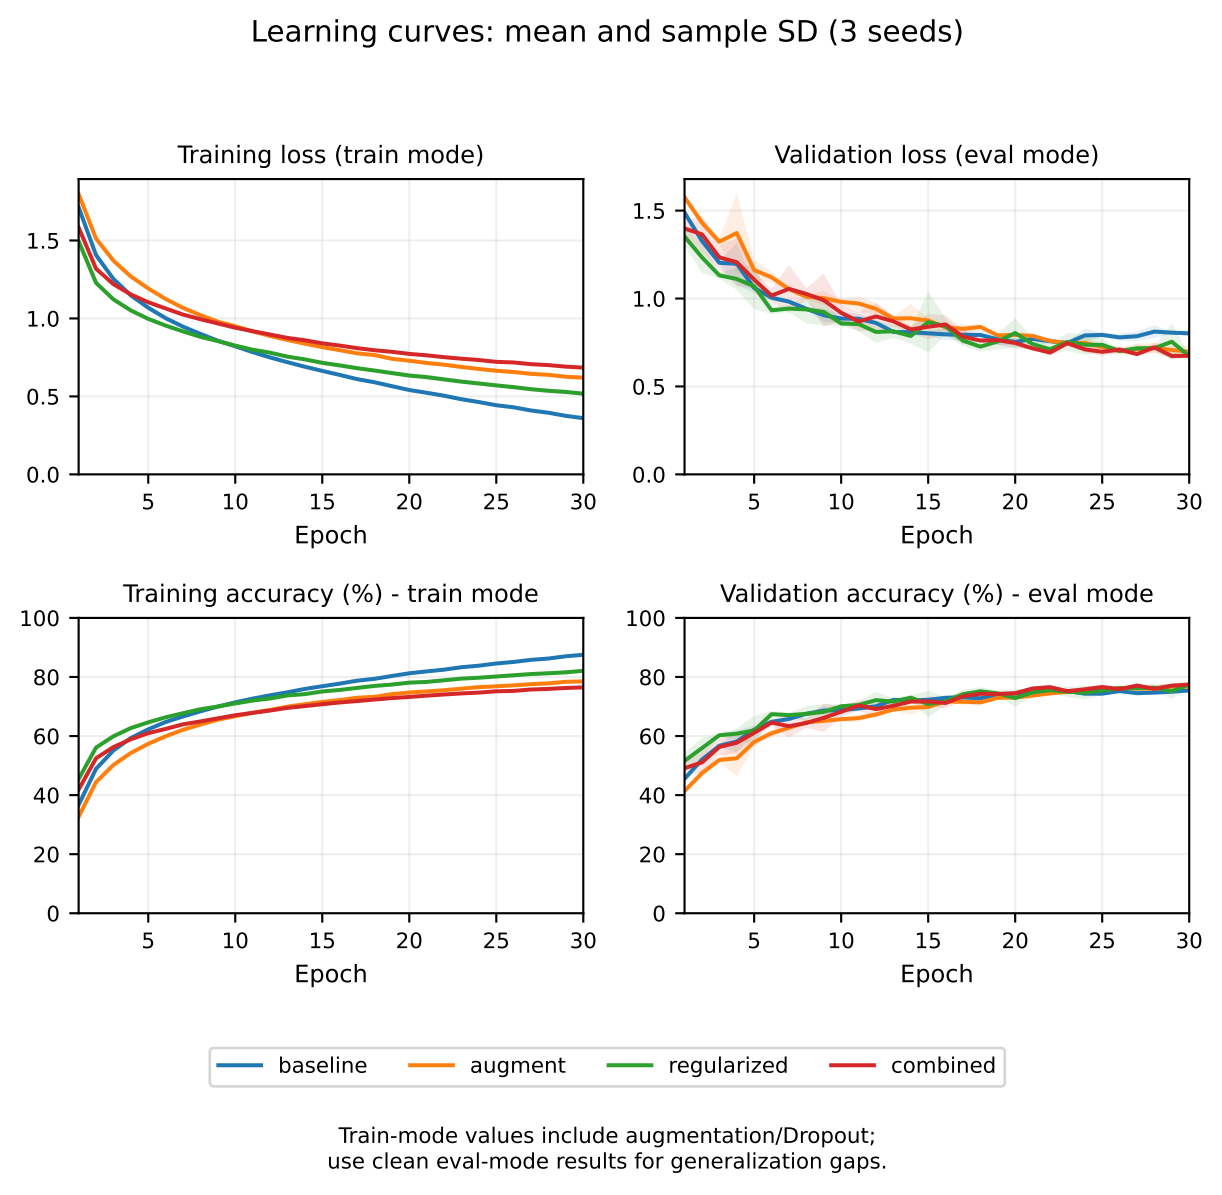
Learning curves: mean and sample SD (3 seeds)
Training loss (train mode)
Validation loss (eval mode)
Training accuracy (%) - train mode
Validation accuracy (%) - eval mode
Epoch
baseline
augment
regularized
combined
Train-mode values include augmentation/Dropout;
use clean eval-mode results for generalization gaps.